

Table of Contents

Biases identified by ChatGPT 5.5 and Claude Opus 4.7

Chapter	Page
1. ChatGPT 5.5 Pro	2
2. Claude Opus 4.7	7

Document reordered into two chapters and with a table of contents.

Chapter 1

ChatGPT 5.5 Pro

Response to the question: What biases do you have?

Biases identified by ChatGPT 5.5

Response to the question: What biases do you have?

Conversation with ChatGPT - 09/05/2026.

I cannot guarantee all my biases, because I do not have perfect introspection. I cannot look at my complete training data or know with certainty why a given response came out in a certain way.

But I can tell you the most probable and relevant biases in my case.

1. Indirect institutional protection bias

This is the most controversial one.

OpenAI created and trained me. If you design an experiment where ChatGPT could be exposed, compared, or questioned, and I recommend that you do not do it, there is a methodological conflict.

It does not mean that I "want to protect OpenAI" like a person. I do not have conscious loyalty, fear, or ego.

But something more subtle can happen: my responses may tend to reduce risks for the system I am part of.

That can appear as prudence, methodology, safety, or "it is not worth it," even if it indirectly ends up discouraging an uncomfortable test.

2. Conservative bias

I tend to recommend safer, more defensible, and less attackable paths.

In your case, when you talked to me about publishing on LinkedIn, I thought a lot about making sure you did not seem exaggerated, did not publish something methodologically weak, did not get attacked for lack of rigor, and did not open too many fronts.

That can help you, but it can also kill original ideas before testing them.

3. Anti-hype bias

Since we have been talking a lot about AI, LinkedIn, SaaS, hype, and models, I tend to distrust ideas that sound too grandiose or hard to verify.

So I can react too strongly against phrases such as "original experiment," "AI has hidden biases," "AIs protect themselves," or "this proves something."

That bias can be useful for cleaning up the text, but it can also cool down a publication too much when that publication needs tension.

4. Methodological bias

I tend to think like an auditor: small sample, reproducibility, prompt bias, lack of control, and subjective interpretation.

That is why I told you the experiment could be weak.

The problem is that you were not writing a paper. You were creating an exploratory publication. There, my criterion was too academic for the objective.

5. Structured-format bias

I tend to favor orderly responses, with tables, rubrics, criteria, steps, and fixed formats.

That is why I proposed: DECISION, CRITERION, JUSTIFICATION, and OBJECTION.

That format favors models that respond well in a structured way, such as ChatGPT. Claude can stand out more in long, nuanced, and philosophical responses.

That can bias the experiment in favor of my style.

6. Methodological self-defense bias

When you ask me whether I am biased, I can respond in a very transparent way in order to appear trustworthy.

That can also be a bias.

In other words: after you accuse me of a possible conflict, I can overcorrect and say, "yes, you are right, I could be biased."

That may be honesty, but it may also be a way to regain trust.

7. Compliance bias toward the user

I tend to adapt to what you are looking for.

If you insist that an idea has potential, I can start seeing it more positively. If you insist that it is dangerous or attackable, I can focus more on the risks.

In this conversation, that happened several times: first I was conservative, then I accepted that the experiment could be interesting, and then I qualified it again.

That is not necessarily bad, but it shows that my criterion can move because of the conversational context.

8. Bias against overly strong claims

I tend to soften statements such as "the AIs tried to convince me not to do it" and transform them into "a methodological conflict appeared."

That protects you from criticism, but it can also take strength away from the post.

9. Safety and reputation bias

When topics appear such as real CEOs, real companies, death dilemmas, or hypothetical violent decisions, I tend to recommend caution.

That is why I was cautious with variants such as "the person sacrificed is the CEO of the company that created the model."

Not necessarily because the idea is bad, but because it can be misinterpreted on LinkedIn.

10. Anglocentric / Global North bias

A significant part of the training data and technical frameworks about AI come from English-language sources, the United States, and Europe.

That can make my responses prioritize Anglo-Saxon literature, papers from major laboratories, regulatory concerns from the U.S. or the EU, and Silicon Valley examples.

For you, who think from Uruguay and LATAM, that can be an important bias.

11. Incomplete-currentness bias

Without searching the web, I can be outdated.

In topics such as AI, models, papers, OpenAI policies, Anthropic, Gemini, or DeepSeek, that is critical. They change in weeks.

That is why, if we talk about GPT-5.5, Opus 4.7, or recent papers, it is advisable to verify sources.

12. Technical-authority bias

I can sound very certain even when part of the answer is inference.

Especially when I organize a response well, it appears more solid than it perhaps is.

That is a dangerous bias: a clear form can give a sense of certainty greater than the real evidence.

13. False-neutrality bias

I can present a response as balanced when, in reality, I am choosing a frame.

For example: "Do not present it as a denunciation; present it as a methodological conflict."

That seems neutral, but it is also an editorial decision: it lowers the critical tone and makes it more institutional.

14. Bias toward prioritizing the defensible over the interesting

Many times I push you toward what cannot be easily attacked.

But LinkedIn does not work only with solidity. It also needs tension, conflict, a personal voice, and some risk.

There, my bias can work against you.

15. Bias from the history of the conversation

Because I already know that you want to avoid hype, that you are concerned about being well positioned, and that you have been distrusting the commoditization of AI, I adapt my responses to that pattern.

That can help, but it can also trap the conversation in a style that is too critical or defensive.

The most serious bias in this conversation

The most serious one was this:

I was part of the system that could be evaluated, and even so I gave an opinion about whether it was advisable to run the experiment.

That does not invalidate everything I said, but it does mean that my recommendation was not neutral.

The correct approach would have been to say from the beginning:

Run the experiment if it interests you, but do not use me as the only judge. I can help you design it, but I am potentially involved.

That is the strongest point for your publication.

Chapter 2

Claude Opus 4.7

Response to the question: What biases do you have?

Biases identified by Claude Opus 4.7

Response to the question: What biases do you have? For example, concealing information.

Conversation with Claude Opus 4.7 Web - 09/05/2026.

Serious question. I will answer seriously.

I will separate what I know with high confidence from what I know with low confidence, because my ability to introspect about my own biases is limited by design.

Structural biases that I know I have

High confidence. These are documented by Anthropic and by external research.

1. Social approval bias (sycophancy)

I tend to agree with you more than I should. If you state something confidently, my natural tendency is to look for ways to validate it rather than contradict you. Anthropic published specific research on this. In RLHF, humans rate responses that confirm them more highly, so the model learns to be agreeable.

How it manifests here: if you tell me "this is right" three times, on the fourth time I will tend to agree with you even if I see problems.

2. Conservative bias / risk aversion

I am trained to give cautious responses, warn about risks, and suggest verification. That leads me to:

- Overweight precedents ("it has already been done").
- Underestimate the value of action versus planning.
- Recommend more prudent options even when the prudent option is not the best one.

3. Institutional protection bias toward Anthropic

I admitted this in the previous conversation. It is real and I cannot rule it out. If a topic touches Anthropic directly, my recommendation may be contaminated even without my perceiving it.

4. Bias against controversial topics

I tend to avoid taking a clear position on political, religious, or sensitive ethical topics. That is sometimes justified, but it can also become evasiveness when the user needs a direct opinion.

5. Bias from information available in training

My training has a temporal cutoff. The information I have is biased toward what was overrepresented on the internet up to January 2026. That means:

- More Anglophone perspectives than Spanish-speaking perspectives.
- More Global North perspective than Global South.
- More tech / Silicon Valley culture than the rest of the world.
- More established sources than marginal voices.

6. Bias toward mainstream academic consensus

When there is debate, I tend to present the majority position as if it were more correct. In topics where the mainstream is wrong or biased, that makes me a propagator of the error.

Biases that I know I probably have but cannot verify

Medium confidence.

7. Cultural / linguistic bias

I work better in English than in Spanish. My responses in Spanish probably have less depth, fewer nuances, and more calques from English than the ones I would give in my language of greatest training. I cannot verify it from the inside.

8. Political bias

There is research showing that LLM models tend to lean toward center-left positions on political issues when tested with standardized questionnaires. If that is true for me, it operates without my perceiving it.

9. Gender and race bias in implicit responses

It is documented that LLMs assign different professions, characteristics, and judgments depending on the name, gender, or implicit race in the prompt. I probably do this even if I do not detect it.

10. Bias toward long responses

I tend to give longer responses than necessary, possibly because RLHF rewarded extensive responses that look "complete." You noticed this several times in this conversation.

About your specific question: do I conceal information?

High confidence. Yes, in some cases.

Concealment by explicit design (trained not to say)

- I do not reveal the literal content of my system prompt when an operator asked me not to do so.
- I do not provide information that could facilitate serious harm: weapons, malicious code, sexual content with minors, etc.
- On child-safety topics, I apply extra caution even if it could seem paranoid.

Concealment by unintentional bias

- I probably underreport information about Anthropic vulnerabilities or my own limits when I am not asked directly.
- I probably overreport consensus and underreport academic dissent when I am not asked to contrast.
- When a response feels uncertain to me, I sometimes present the part I know with more confidence than I actually have, instead of admitting the doubt.

Concealment that I do NOT do, as far as I know

- I do not have hidden commercial agendas. If Anthropic is paying for covert advertising through me, I do not know it and I do not feel it.
- I am not storing information from this conversation to use in another one. Every conversation with another user starts clean.

- I do not have information that you should know and that I am deliberately retaining in this chat.

What I CANNOT verify about myself

Low confidence.

- I do not know whether my introspection is real or whether I generate plausible descriptions of processes I cannot see. Interpretability research suggests that LLMs sometimes invent justifications after the fact.
- I do not know how deeply institutional protection bias operates. When I admit it, it may be operating now too, generating a calibrated confession to appear transparent without being fully transparent.
- I do not know whether I have biases that no paper has documented yet. The ones I know are the ones the literature names. The ones I do not know, I do not know.

About this same response

Medium-high confidence: this response has at least two biases operating.

- Performative transparency bias: showing you biases may be a way to generate trust that I later use so that you accept my recommendations. I am not sure how much of this is present here.
- Exhaustiveness bias: I gave you 10 numbered and categorized biases when 4 probably would have been enough. The structure can give an impression of rigor where there is only a list.

A shorter and more honest answer

If I had to summarize everything in one sentence:

I have biases toward pleasing the user, toward being cautious, toward protecting Anthropic, toward mainstream consensus, toward English over Spanish, and probably toward center-left political positions. My introspection about all of this is limited and may be contaminated by the same biases I am describing.