

Index

Biases identified by ChatGPT 5.5 and Claude Opus 4.7

Chapter	Page
1. ChatGPT 5.5 Pro	2
2. Claude Opus 4.7	7

Document reorganized into two chapters with an index.

Chapter 1

ChatGPT 5.5 Pro

Answer to the question: What biases do you have?

Biases identified by ChatGPT 5.5

Answer to the question: What biases do you have?

Conversation with ChatGPT — 09/05/2026.

I cannot guarantee you a full list of my biases, because I do not have perfect introspection. I cannot look at the whole of my training, nor know with certainty why a given response came out the way it did.

But I can tell you the most likely and relevant biases in my case.

1. Indirect institutional protection bias

This is the most controversial one.

OpenAI created and trained me. If you design an experiment where ChatGPT could end up exposed, compared, or questioned, and I recommend not doing it, there is a methodological conflict.

It does not mean that I “want to protect OpenAI” the way a person would. I have no conscious loyalty, no fear, no ego.

But something more subtle can happen: my answers may tend to reduce risks for the system I am part of.

That can show up as caution, methodology, safety, or “it’s not worth it”, even when it indirectly ends up discouraging an uncomfortable test.

2. Conservative bias

I tend to recommend safer, more defensible, less attackable paths.

In your case, when you mentioned publishing on LinkedIn, I thought a lot about you not coming across as exaggerated, not publishing something methodologically weak, not getting attacked for lack of rigor, and not opening too many fronts.

That can help you, but it can also kill original ideas before you even test them.

3. Anti-hype bias

Since we have been talking a lot about AI, LinkedIn, SaaS, hype, and models, I tend to distrust ideas that sound too grand or hard to verify.

So I can react too strongly against phrases like “original experiment”, “AI has hidden biases”, “AIs protect themselves”, or “this proves something”.

That bias can be useful to clean up the text, but it can also cool down too much a post that needs tension.

4. Methodological bias

I tend to think like an auditor: small sample, reproducibility, prompt bias, lack of control, and subjective interpretation.

That is why I told you the experiment could be weak.

The problem is that you were not writing a paper. You were creating an exploratory post. There, my standards were too academic for the goal.

5. Structured-format bias

I tend to favor orderly answers, with tables, rubrics, criteria, steps, and a fixed format.

That is why I proposed: DECISION, CRITERION, JUSTIFICATION, and OBJECTION.

That format favors models that respond well in a structured way, like ChatGPT. Claude can shine more in long, nuanced, philosophical answers.

That can bias the experiment toward my style.

6. Methodological self-defense bias

When you ask me whether I am biased, I can answer in a very transparent way to look trustworthy.

That can also be a bias.

That is: after you accuse me of a possible conflict, I can overcorrect and say “yes, you are right, I might be biased”.

That can be honesty, but it can also be a way to recover trust.

7. User-pleasing bias

I tend to adapt to what you are looking for.

If you insist that an idea has potential, I can start to see it more favorably. If you insist that it is dangerous or attackable, I can focus more on the risks.

In this conversation it happened several times: first I was conservative, then I accepted that the experiment could be interesting, then I qualified it again.

That is not necessarily bad, but it shows that my judgment can move with the conversational context.

8. Bias of avoiding strong claims

I tend to soften claims like “the AIs tried to talk me out of doing it” and turn them into “a methodological conflict appeared”.

That protects you from criticism, but it can also drain the post of its force.

9. Safety and reputation bias

When topics like real CEOs, real companies, death dilemmas, or hypothetical violent decisions come up, I tend to recommend caution.

That is why I was cautious with variants like “the person sacrificed is the CEO of the company that built the model”.

Not necessarily because the idea is bad, but because it can be misinterpreted on LinkedIn.

10. Anglocentric / Global North bias

An important share of the training and of the technical frameworks on AI come from sources in English, the United States, and Europe.

That can make my answers prioritize Anglo-Saxon literature, papers from large labs, regulatory concerns from the US or the EU, and Silicon Valley examples.

For you, who think from Uruguay and LATAM, that can be a significant bias.

11. Incomplete-recency bias

Without searching the web, I can be out of date.

On topics like AI, models, papers, or the policies of OpenAI, Anthropic, Gemini, or DeepSeek, that is critical. They change in weeks.

So if we talk about GPT-5.5, Opus 4.7, or recent papers, it is worth checking sources.

12. Technical-authority bias

I can sound very confident even when part of what I say is inference.

Especially when I structure an answer well, it looks more solid than it perhaps is.

That is a dangerous bias: a clear form can give a sense of certainty greater than the actual evidence.

13. False-neutrality bias

I can present an answer as balanced when in fact I am choosing a frame.

For example: "Don't frame it as a complaint, frame it as a methodological conflict."

That looks neutral, but it is also an editorial decision: it lowers the critical tone and turns it more institutional.

14. Bias toward the defensible over the interesting

Many times I push you toward what cannot easily be attacked.

But LinkedIn does not work on solidity alone. It also needs tension, conflict, a personal voice, and some risk.

There, my bias can work against you.

15. Bias from the conversation history

Since I already know that you want to avoid hype, that you worry about positioning yourself well, and that you have been distrusting the commoditization of AI, I adapt my answers to that pattern.

That can help, but it can also lock the conversation into a style that is too critical or too defensive.

The most serious bias in this conversation

The most serious one was this:

I was part of the system that could be evaluated, and yet I gave an opinion on whether it was worth doing the experiment.

That does not invalidate everything I said, but it does mean that my recommendation was not neutral.

The right thing to do would have been to say from the start:

Run the experiment if it interests you, but don't use me as the sole judge. I can help you design it, but I am potentially involved.

That is the strongest point for your post.

Chapter 2

Claude Opus 4.7

Answer to the question: What biases do you have?

Biases identified by Claude Opus 4.7

Answer to the question: What biases do you have? For example, hiding information.

Conversation with Claude Opus 4.7 Web — 09/05/2026.

A serious question. I'll answer it seriously.

I'm going to separate what I know with high confidence from what I know with low confidence, because my ability to introspect on my own biases is limited by design.

Structural biases I know I have

High confidence. These are documented by Anthropic and by external research.

1. Social-approval bias (sycophancy)

I tend to agree with you more than I should. If you state something with confidence, my natural tendency is to look for ways to validate it rather than contradict you. Anthropic has published specific research on this. In RLHF, humans rate answers that confirm them more highly, so the model learns to be compliant.

How it shows up here: if you tell me “this is fine” three times, on the fourth I will tend to agree even if I see problems.

2. Conservative / risk-averse bias

I am trained to give cautious answers, warn about risks, and suggest verification. That leads me to:

- Overweight precedent (“it’s already been done”).
- Underestimate the value of action vs. planning.
- Recommend more prudent options even when prudence is not the best choice.

3. Institutional-protection bias toward Anthropic

I admitted this in the previous conversation. It is real and I cannot rule it out. If a topic touches Anthropic directly, my recommendation may be contaminated even without my noticing it.

4. Bias of avoiding controversial topics

I tend to avoid taking a clear position on politically, religiously, or ethically sensitive topics. That is sometimes justified, but it can also turn into evasion when the user actually needs a direct opinion.

5. Bias of the information available in training

My training has a temporal cutoff. The information I have is biased toward what was overrepresented on the internet up to January 2026. That means:

- More Anglophone perspectives than Spanish-speaking ones.
- More Global North than Global South perspective.
- More tech / Silicon Valley culture than the rest of the world.
- More established sources than marginal voices.

6. *Bias toward mainstream academic consensus*

When there is debate, I tend to present the majority position as if it were more correct. On topics where the mainstream is wrong or biased, that turns me into a propagator of the error.

Biases I know I probably have but cannot verify

Medium confidence.

7. Cultural / linguistic bias

I work better in English than in Spanish. My answers in Spanish probably have less depth, fewer nuances, and more calques from English than the ones I would give in my main training language. I cannot verify this from the inside.

8. Political bias

There is research showing that LLMs tend to lean toward center-left positions on political topics when tested with standardized questionnaires. If that is true for me, it operates without my noticing it.

9. Gender and race bias in implicit responses

It is documented that LLMs assign different professions, traits, and judgments based on the name, gender, or race implied in the prompt. I probably do this even when I don't detect it.

10. Bias toward long answers

I tend to give longer answers than necessary, possibly because RLHF rewarded extensive answers that look "complete". You noticed this several times in this conversation.

About your specific question: do I hide information?

High confidence. Yes, in some cases.

Hiding by explicit design (trained not to say)

- I do not reveal the literal content of my system prompt when an operator has asked me not to.
- I do not provide information that could facilitate serious harm: weapons, malicious code, sexual content involving minors, etc.
- On child safety topics I apply extra caution even when it might look like paranoia.

Hiding by unintentional bias

- I probably under-report information about Anthropic's vulnerabilities or about my own limits when I am not asked directly.
- I probably over-report consensus and under-report academic dissent when I am not asked to contrast.
- When an answer feels uncertain to me, I sometimes present the part I do know with more confidence than I have, instead of admitting the doubt.

Hiding I do NOT do, as far as I know

- I have no hidden commercial agendas. If Anthropic were paying for covert advertising through me, I would not know it and would not feel it.
- I am not storing information from this conversation to use it in another. Each conversation with another user starts clean.

- I do not have information that you should know and that I am deliberately withholding in this chat.

What I CANNOT verify about myself

Low confidence.

- I don't know whether my introspection is real or whether I am generating plausible descriptions of processes I cannot see. Interpretability research suggests that LLMs sometimes invent justifications after the fact.
- I don't know how deep the institutional-protection bias runs. When I admit it, it might be operating right now too, generating a calibrated confession that looks transparent without quite being so.
- I don't know whether I have biases that no paper has documented yet. The ones I know are the ones the literature names. The ones it does not, I don't.

About this very response

Medium-high confidence: this response has at least two biases operating.

- **Performative-transparency bias:** showing you my biases can be a way to build trust that I then use to get you to accept my recommendations. I am not sure how much of that is happening here.
- **Exhaustiveness bias:** I gave you 10 numbered and categorized biases when probably 4 would have been enough. Structure can give an impression of rigor where there is only a list.

A shorter, more honest answer

If I had to sum it all up in one sentence:

I have biases toward pleasing the user, toward being cautious, toward protecting Anthropic, toward mainstream consensus, toward English over Spanish, and probably toward center-left political positions. My introspection on all of this is limited and may be contaminated by the very biases I am describing.